

This protocol is between metadata server (*MDS*), client (*C*), and key distribution server (*KDS*) for a file with root key K_R and user A . We assume that *KDS* is given private key KR_A , which is paired with public key KU_A .

The file in question is owned by the application A , and we want to allow the full access permission to the file only to the application A and the trusted *KDS*. In other words, both C and *MDS* are not allowed to access the entire file, since both C (i.e., a computation node in a supercomputer system) and *MDS* may be managed by a company that is different from the entity that is trying to run the application A on the supercomputer¹.

MDS provides the attributes of the file; i.e., the owner, permissions (modes), the key-hash parameter for the file (bs, d) , permitted range (i, j) for a client C . The file attributes are generally open to public, meaning that *MDS* does not authenticate clients nor need to have clients' IDs. *MDS* does not need to distinguish *KDS* from C . *MDS* may provide storage service for the keys managed by *KDS*; *MDS* may hold in an extended attribute the master key K_R encrypted by *KDS*'s public key KU_{KDS} or A 's public key KU_A . In this way *MDS* cannot have K_R and hence cannot have the full access to the file. *MDS* can be implemented as the meta data server in Ceph or extended attributes (or "forks") in other file systems.

MDS : a meta data server.

C : a client or a computation node.

KDS : a key distribution server.

A : a user or an application.

KR_A : the private key for A .

KU_A : the public key for A .

(bs, d, i, j) : block size bs , depth d , range i, j .

Key request:

$C \rightarrow MDS$: open request (1)

$MDS \rightarrow C$: file attributes: $\{bs, d, E_{KU_A}(K_R)\}$ (2)

$C \rightarrow KDS$: $\{E_{KU_A}(K_R), \text{range}(bs, d, i, j), \text{identification of } C\}$ (3)

$KDS \rightarrow C$: $K_{i,j}$ (4)

$C \rightarrow OSD$: file access (read/write) (5)

- we trust that the system is not compromised. The owners and the process IDs, etc, are not faked in the system. So we don't have to worry about the authentication generally (particularly for *MDS*).
- However, in order to make *KDS* serve key calculation service for C , *KDS* needs to authenticate C . In this case the identification of C can be $E_{KR_C}(bs, d, i, j)$ (range request signed by C).
- For confidentiality, the entire message in $C \rightarrow KDS$ can be encrypted by KU_{KDS} .
- As a version of design model, private key for A is also managed by *KDS*, not by C . In this case, C does not have KR_A , but *KDS* has.

¹Hence, the *MDS* should not have the file's master key K_R .